

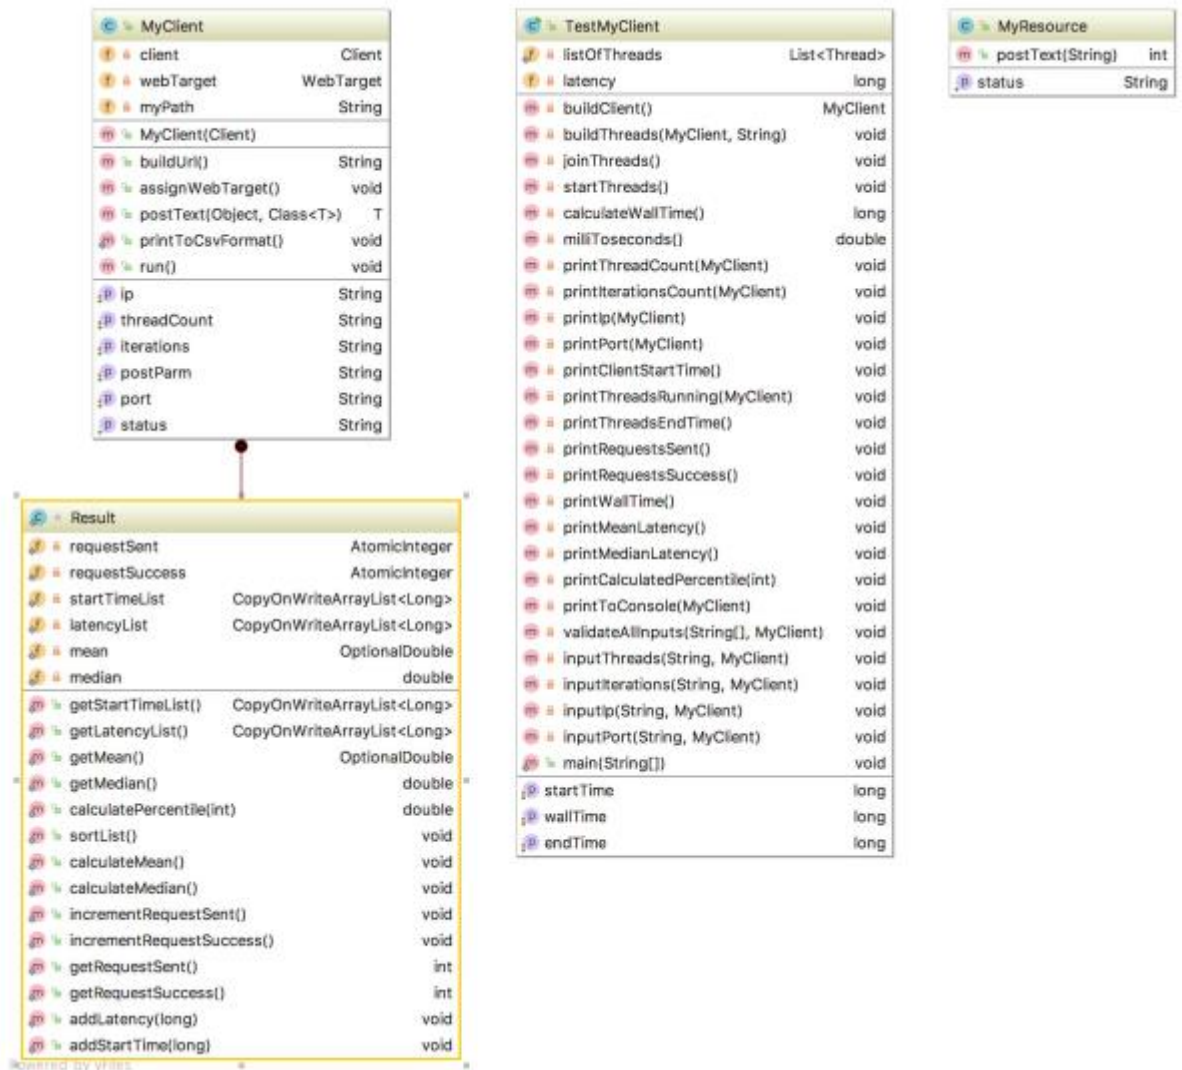
Building Scalable Distributed Systems – CS6650

Assignment – 1

<https://github.com/kinshukjuneja/BSDS-CS6650>

Name: Kinshuk Juneja

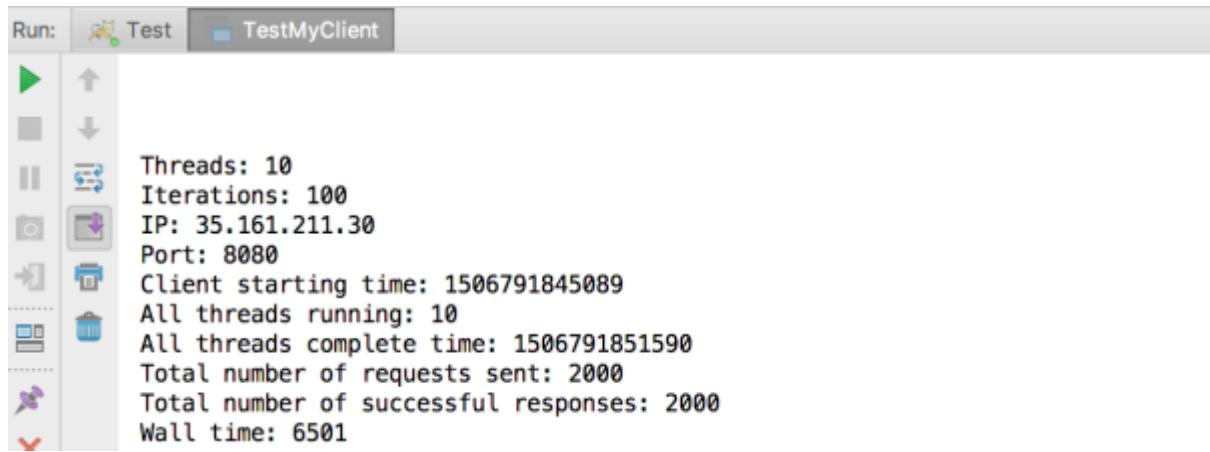
- 1) A 1 page overview of your design (a simple block diagram would suffice). The aim is to quickly summarize your design so emphasize important components and abstractions. (1 point)
 - I have 4 classes in total where Result is an inner class of MyClient which in turn implements Runnable.
 - On execution of TestMyClient containing main() method it first builds a client and then creates a list of all threads taken from input or default to 10 and joins those thread so that code following main() is not executed. On starting all the threads, Client's overridden run() method is executed where requests for GET/POST till number of iterations(default: 100) is run for each GET and POST and returns the value from MyResource class.
 - The inner class Result is helpful to calculate number of requests sent, number of successful requests, wall time, mean, median, percentile calculation, so that TestMyClient class can display those in output.
 - To plot graphs, I also included list of start time for each request in GET/POST along with latency which is calculated as (end time of request – start time of request).
 - Please refer to the block diagram below for an overview:



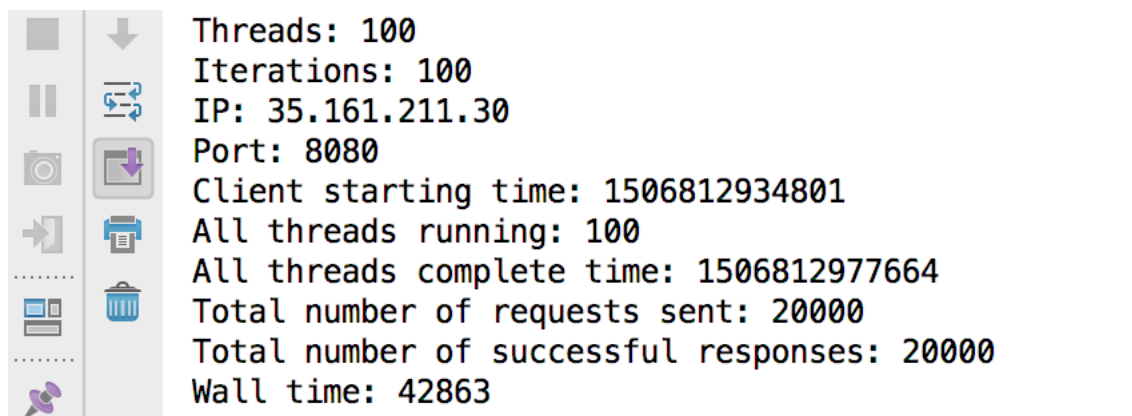
- 2) URL to your git repo. (3 points)
- <https://github.com/kinshukjuneja/BSDS-CS6650>

3) Two screenshots for step 4 showing correct execution and completion of the two specified tests (**3 points for each**)

➤ **GET/POST – 10/100:**

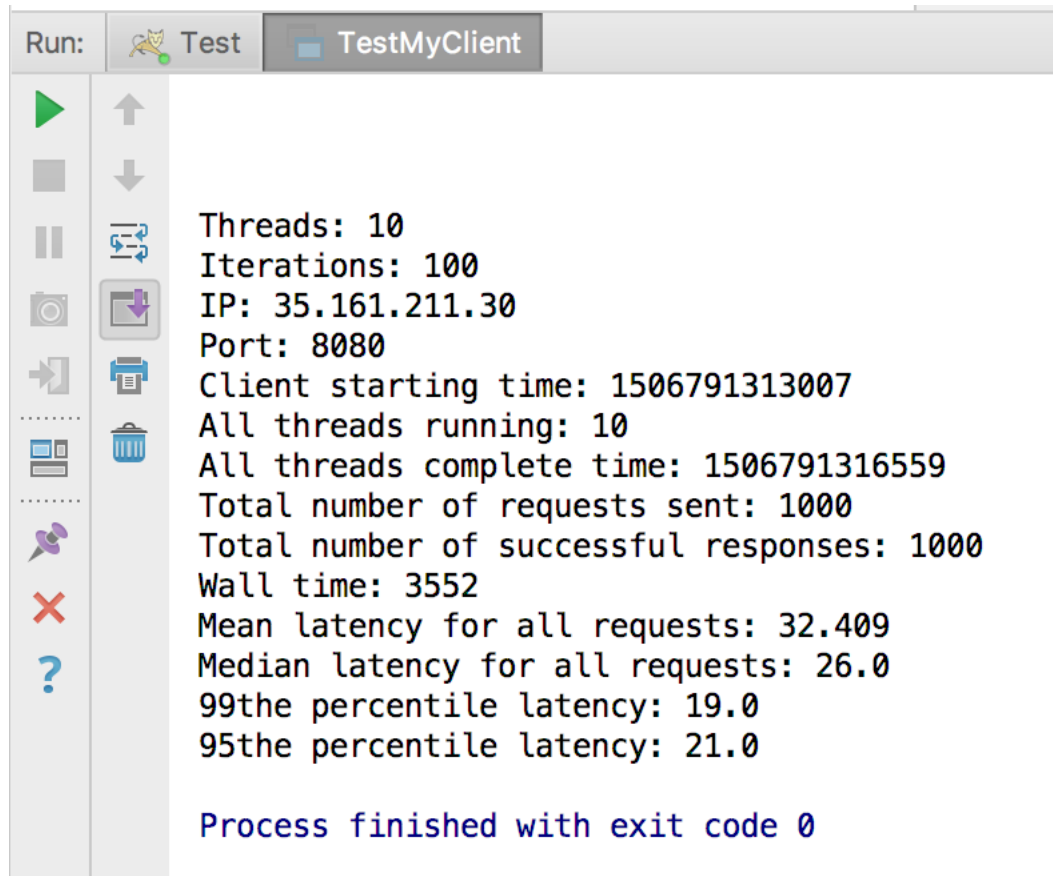


➤ **GET/POST – 100/100**

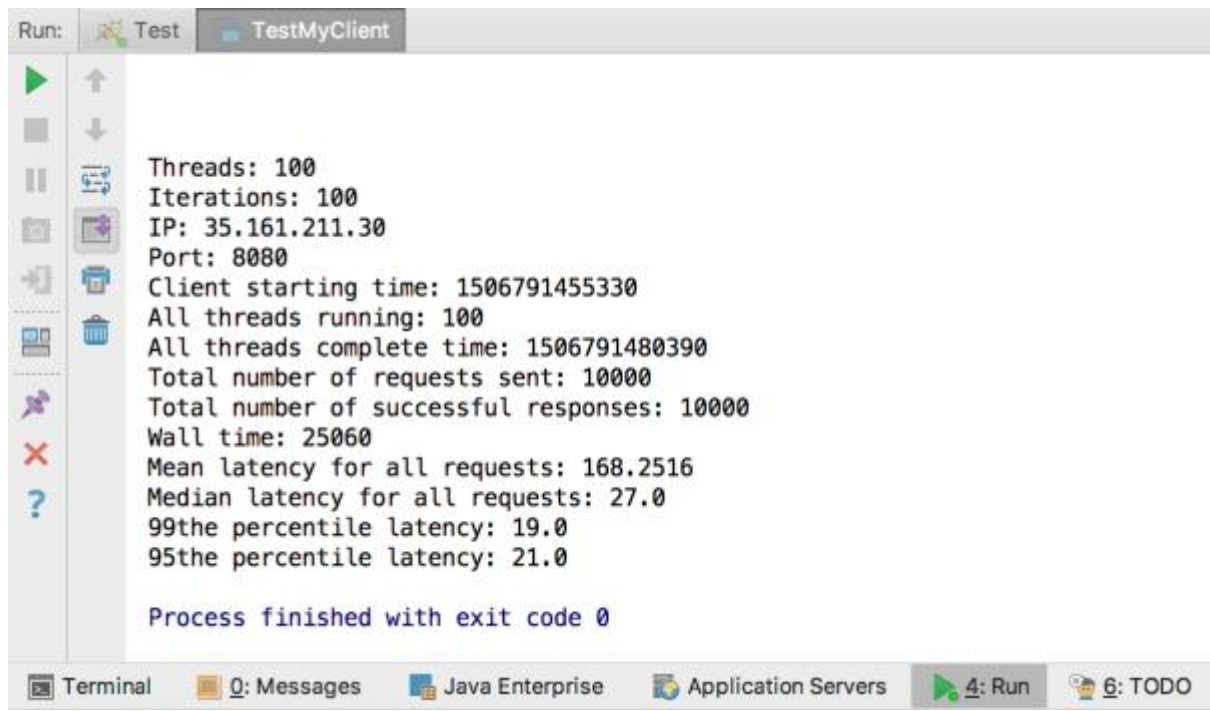


4) Two screenshots for step 5 showing correct execution and completion of the two specified tests. If you use an additional tool like a spreadsheet, show the results in this in addition to the two screenshots showing the test running (**5 points for each**)

➤ **GET – 10/100:**



➤ GET – 100/100

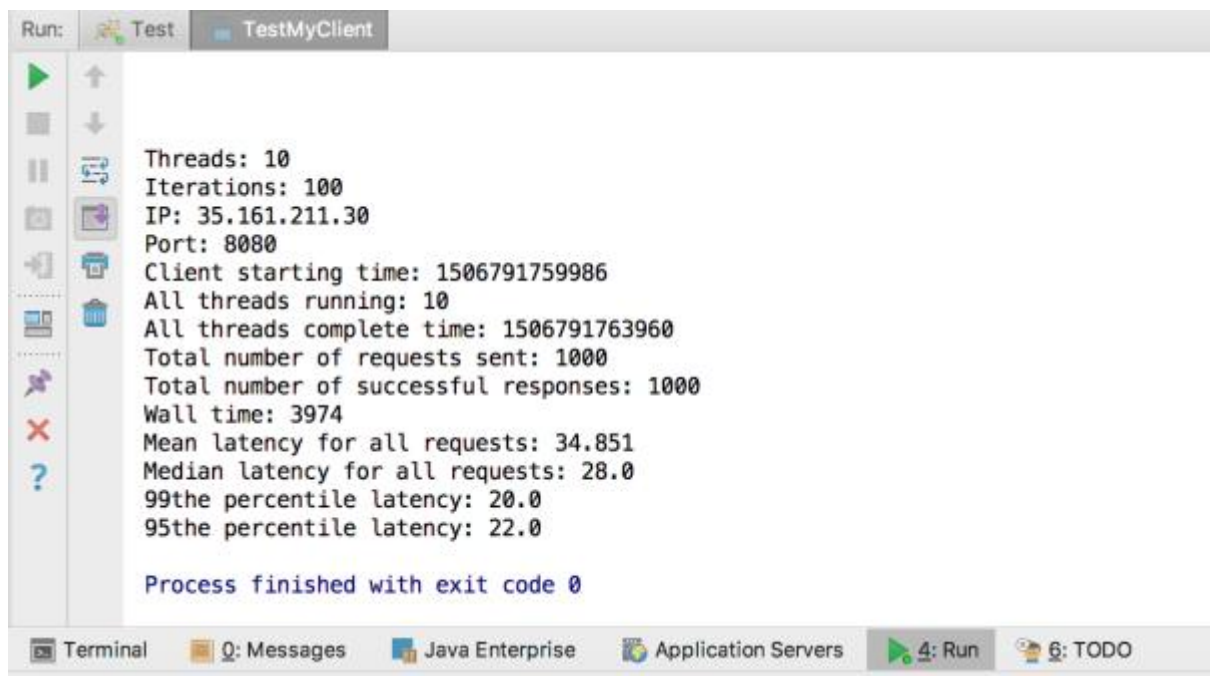


```
Run: Test TestMyClient

Threads: 100
Iterations: 100
IP: 35.161.211.30
Port: 8080
Client starting time: 1506791455330
All threads running: 100
All threads complete time: 1506791480390
Total number of requests sent: 10000
Total number of successful responses: 10000
Wall time: 25060
Mean latency for all requests: 168.2516
Median latency for all requests: 27.0
99the percentile latency: 19.0
95the percentile latency: 21.0

Process finished with exit code 0
```

➤ POST – 10/100

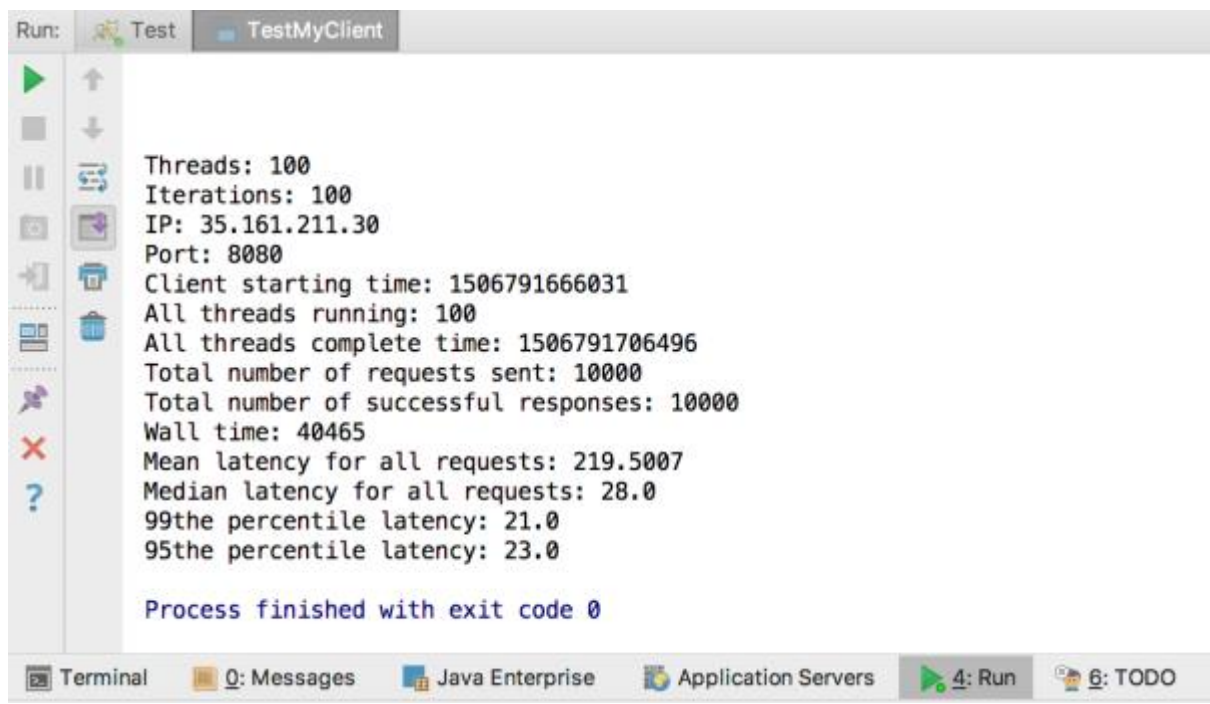


```
Run: Test TestMyClient

Threads: 10
Iterations: 100
IP: 35.161.211.30
Port: 8080
Client starting time: 1506791759986
All threads running: 10
All threads complete time: 1506791763960
Total number of requests sent: 1000
Total number of successful responses: 1000
Wall time: 3974
Mean latency for all requests: 34.851
Median latency for all requests: 28.0
99the percentile latency: 20.0
95the percentile latency: 22.0

Process finished with exit code 0
```

➤ POST – 100/100



The screenshot shows the Run console of an IDE. At the top, the 'Run:' tab is active, showing 'Test' and 'TestMyClient'. The console output displays the following test results:

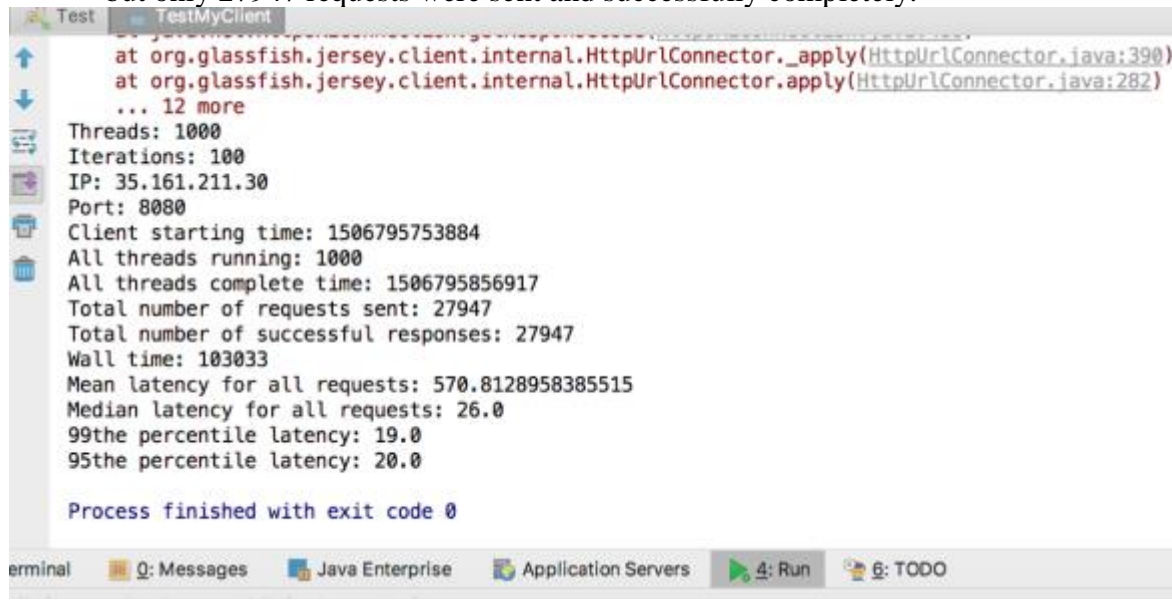
```
Threads: 100
Iterations: 100
IP: 35.161.211.30
Port: 8080
Client starting time: 1506791666031
All threads running: 100
All threads complete time: 1506791706496
Total number of requests sent: 10000
Total number of successful responses: 10000
Wall time: 40465
Mean latency for all requests: 219.5007
Median latency for all requests: 28.0
99th percentile latency: 21.0
95th percentile latency: 23.0

Process finished with exit code 0
```

The bottom of the IDE shows a toolbar with icons for Terminal, Q: Messages, Java Enterprise, Application Servers, 4: Run, and 6: TODO.

5) Step 6 Stress testing: Submit a short (1 page?) report describing what you did to explore the tolerances of your application, what broke it, and how. (1 point)

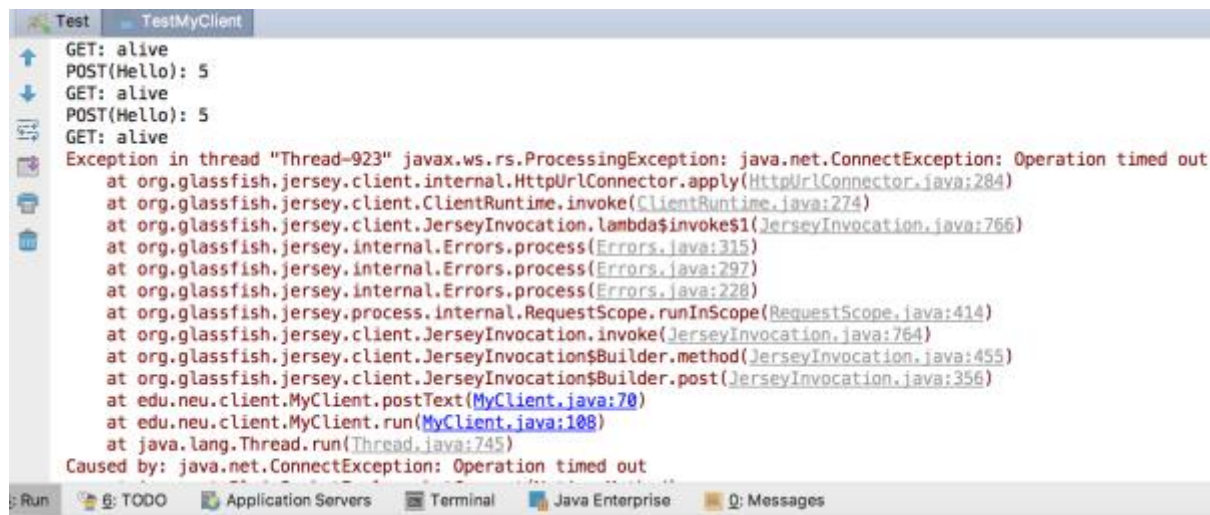
- **Check Tolerance:** To explore the tolerance I gradually increased my number of threads for 100 iterations(requests) and vice-verse. The following result was obtained:
- **What broke it - Get – 1000/100:** On increasing the number of threads to 1000 for 100 iterations(requests) just for GET, the total number of requests should be 100000, but only 27947 requests were sent and successfully completely.



```
Test TestMyClient
at org.glassfish.jersey.client.internal.HttpUrlConnector._apply(HttpUrlConnector.java:390)
at org.glassfish.jersey.client.internal.HttpUrlConnector.apply(HttpUrlConnector.java:282)
... 12 more
Threads: 1000
Iterations: 100
IP: 35.161.211.30
Port: 8080
Client starting time: 1506795753884
All threads running: 1000
All threads complete time: 1506795856917
Total number of requests sent: 27947
Total number of successful responses: 27947
Wall time: 103033
Mean latency for all requests: 570.8128958385515
Median latency for all requests: 26.0
99th percentile latency: 19.0
95th percentile latency: 20.0

Process finished with exit code 0
```

- **How - GET/POST – 1000/100:** On testing for both GET and POST requests simultaneously for 1000 threads and 100 iterations (requests), a total of 200000 requests should have been made, but that did not happen because an error was **first** thrown on making just 17218 requests, after which requests continued but errors appeared in between.
- This was eventually, because in my program **Operation timed out.**
i.e. Caused by: java.net.ConnectException: Operation timed out




```
Test TestMyClient
GET: alive
POST(Hello): 5
GET: alive
POST(Hello): 5
GET: alive
Exception in thread "Thread-923" javax.ws.rs.ProcessingException: java.net.ConnectException: Operation timed out
at org.glassfish.jersey.client.internal.HttpUrlConnector.apply(HttpUrlConnector.java:284)
at org.glassfish.jersey.client.ClientRuntime.invoke(ClientRuntime.java:274)
at org.glassfish.jersey.internal.Errors.process(Errors.java:315)
at org.glassfish.jersey.internal.Errors.process(Errors.java:297)
at org.glassfish.jersey.internal.Errors.process(Errors.java:228)
at org.glassfish.jersey.process.internal.RequestScope.runInScope(RequestScope.java:414)
at org.glassfish.jersey.client.JerseyInvocation.invoke(JerseyInvocation.java:764)
at org.glassfish.jersey.client.JerseyInvocation$Builder.method(JerseyInvocation.java:455)
at org.glassfish.jersey.client.JerseyInvocation$Builder.post(JerseyInvocation.java:356)
at edu.neu.client.MyClient.postText(MyClient.java:70)
at edu.neu.client.MyClient.run(MyClient.java:108)
at java.lang.Thread.run(Thread.java:745)
Caused by: java.net.ConnectException: Operation timed out
```


6) Step 6 charting: Submit a 1 page report detailing your test run, method of calculation, and chart showing latencies. (1 point)

- To plot graphs, I included an ArrayList of start time for each request in GET/POST along with latency which is calculated as (end time of request – start time of request).
- As soon as the threads override Client's run() method, I calculate the start time in milliseconds(System.currentTimeMillis()) and assign it to a variable startTime for each GET/POST tested individually and add this to my list of start times.
- After this, when a GET/POST request is made and completed, I generate an endTime for GET/POST respectively and subtract endTime – startTime to get the latency which is added to an ArrayList of latency.
- For all the graphs listed below: X-axis is my start time of request made and Y-axis is my Latency.


➤ GET – 10/100:



➤ GET – 100/100



➤ **POST – 10/100**



➤ **POST – 100/100**

